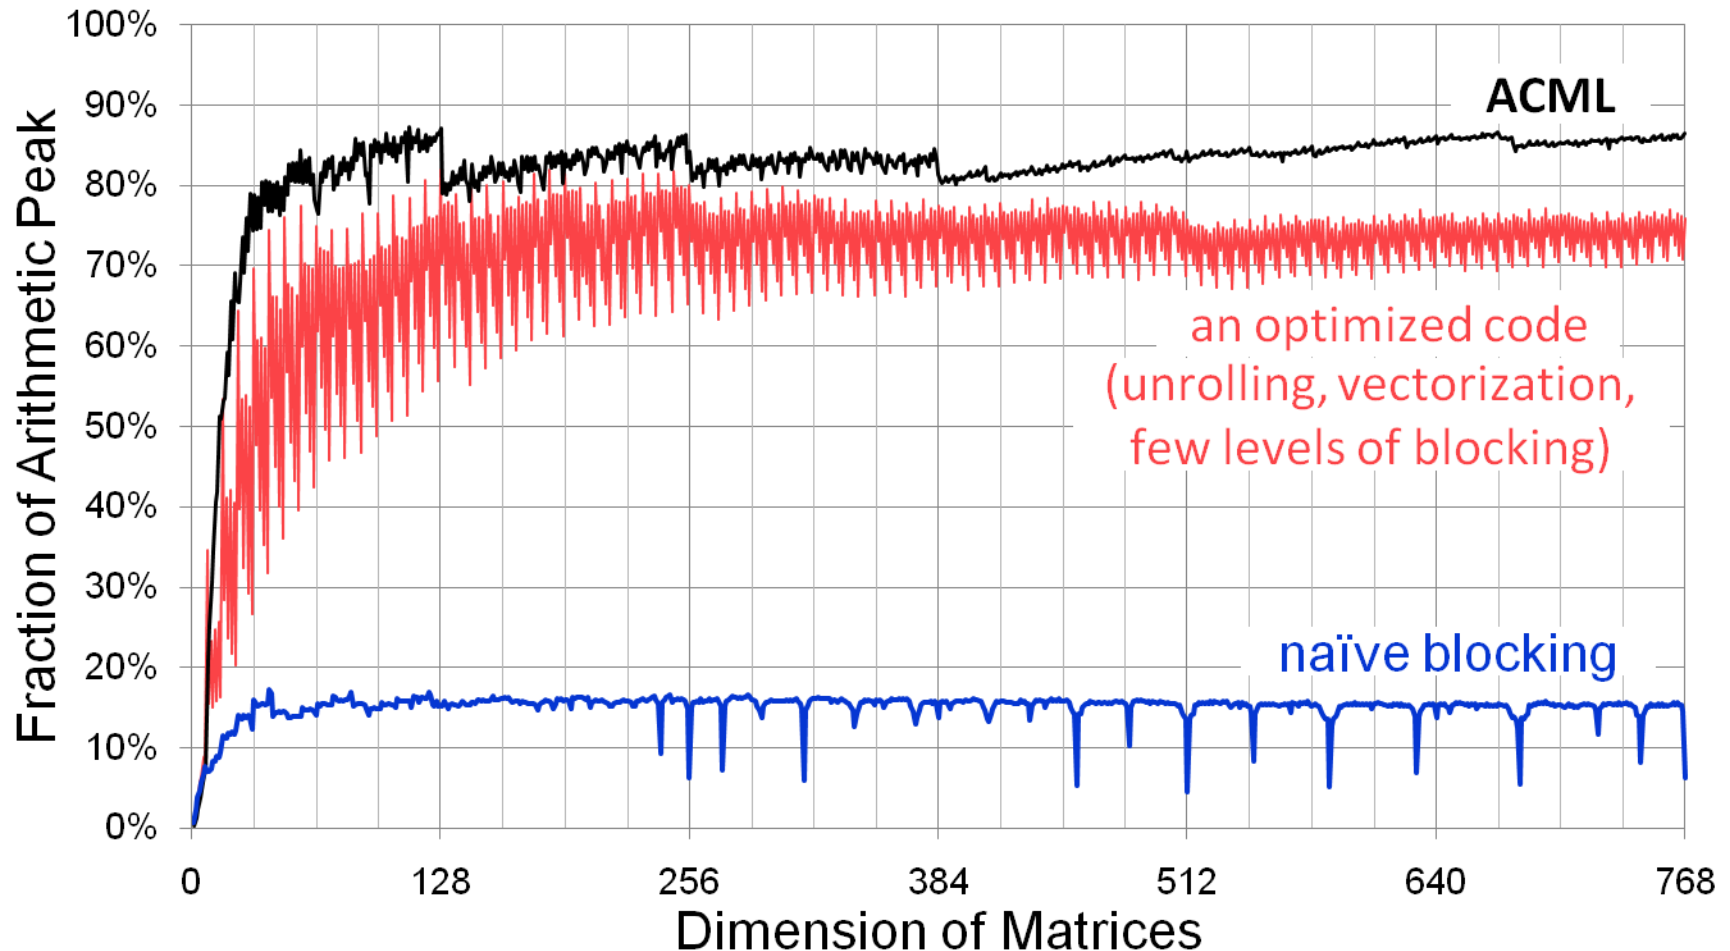


# Notes on Homework 1

What is Peak ?



# Summary of SSE intrinsics

---

**#include <emmintrin.h>**

**Vector data type:**

- **\_\_m128d**

**Load and store operations:**

- **\_mm\_load\_pd**
- **\_mm\_store\_pd**
- **\_mm\_loadu\_pd**
- **\_mm\_storeu\_pd**

**Load and broadcast across vector**

- **\_mm\_load1\_pd**

**Arithmetic:**

- **\_mm\_add\_pd**
- **\_mm\_mul\_pd**

## 2x2 Matrix Multiply

---

$$C_{00} += A_{00}B_{10} + A_{01}B_{00}$$

$$C_{10} += A_{10}B_{10} + A_{11}B_{00}$$

$$C_{01} += A_{00}B_{11} + A_{01}B_{01}$$

$$C_{11} += A_{10}B_{11} + A_{11}B_{01}$$

**Rewrite as SIMD algebra**

$$C00\_C10 += A00\_A10 * B10\_B10$$

$$C01\_C11 += A00\_A10 * B11\_B11$$

$$C00\_C10 += A01\_A11 * B00\_B00$$

$$C01\_C11 += A01\_A11 * B01\_B01$$

## Example: multiplying 2x2 matrices

---

```
#include <emmintrin.h>
c1 = _mm_loadu_pd( C+0*lda );           //load unaligned block in C
c2 = _mm_loadu_pd( C+1*lda );
for( int i = 0; i < 2; i++ )
{
    a = _mm_load_pd( A+i*lda ); //load aligned i-th column of A
    b1 = _mm_load1_pd( B+i+0*lda ); //load i-th row of B
    b2 = _mm_load1_pd( B+i+1*lda );
    c1=_mm_add_pd( c1, _mm_mul_pd( a, b1 ) ); //rank-1 update
    c2=_mm_add_pd( c2, _mm_mul_pd( a, b2 ) );
}
_mm_storeu_pd( C+0*lda, c1 );           //store unaligned block in C
_mm_storeu_pd( C+1*lda, c2 );
```

# Other Issues

---

## Checking efficiency of the compiler helps

- Use -S option to see the generated assembly code
- Inner loop should consist mostly of ADDPD and MULPD ops, ADDSD and MULSD imply scalar computations
- Consider using another compiler
  - Options are PGI, PathScale and GNU
  - Cray C compiler doesn't seem to have any way to use SSE
    - It is capable of auto-vectorizing loops it recognizes
- Look through Goto and van de Geijn's paper
- Don't ignore the other optimizations
  - more compiler flags
  - loop unrolling
  - inlining
  - keywords (ask your classmates!)